

# Notation

To ensure that everyone is up to speed on notation, let's review

- [the notation \(ML Notation.ipynb\)](#) that we used in the "Classical Machine Learning" part of the intro course.
- [additional notation \(Intro to Neural Networks.ipynb\)](#) used in the "Deep Learning" part of the intro course

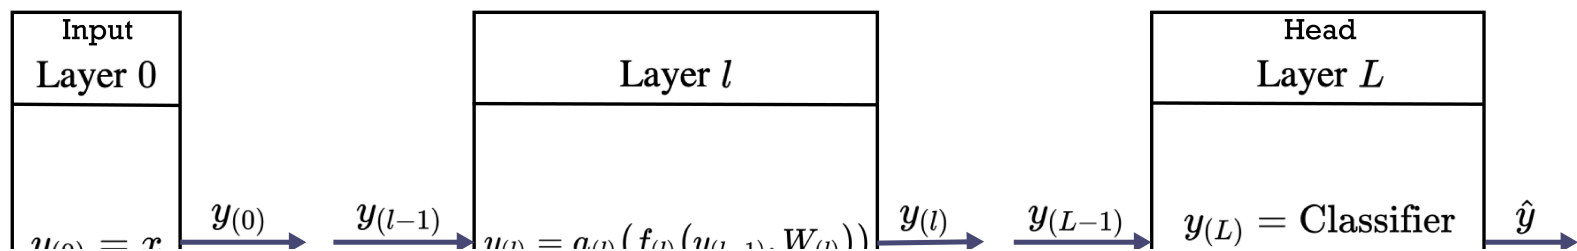
# Representations

A path through a Neural Network can be viewed as a sequence of representation transformations

- transforming *raw features*  $\mathbf{y}_{(0)} = \mathbf{x}$
- into *synthetic features*  $\mathbf{y}_{(l)}$ 
  - varying with layer  $1 \leq l \leq (L - 1)$
- of increasing abstraction

Thus, the output anywhere along the path is an *alternate representation* of the input

**Path through a Neural Network**



Shallow features are less abstract: "syntax", "surface"

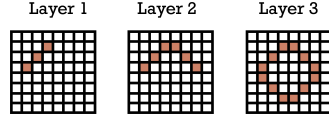
Deeper features are more abstract: "semantics", "concepts"

- We may even interpret the features as "pattern matching" regions or concepts in the raw feature space.

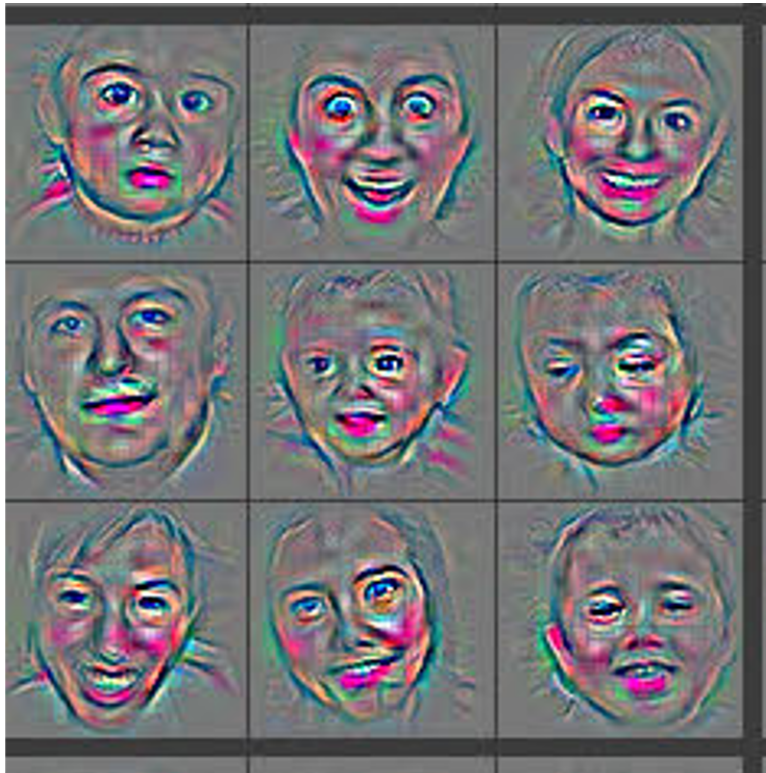
For example, in a CNN

- shallow features are primitive shapes
- deeper features seem to recognize combinations of shallower features

**Input features detected by layer**



Saliency Maps and Corresponding Patches  
Single Layer 5 Feature Map  
On 9 Maximally Activating Input images



Layer 5 ? Feature Map (Row 11, col 1).

Attribution: <https://arxiv.org/abs/1311.2901> (<https://arxiv.org/abs/1311.2901>)

In the simple architectures of the Intro course, we mostly ignored the intermediate representations

$$\mathbf{y}_{(l)} : 1 \leq l \leq (L - 1)$$

The layers were referred to as "hidden" for a reason !

We will discover uses for intermediate representations and show how to build a "feature extractor" to obtain them from a given architecture.

# Recurrent Neural Networks

With a sequence  $\mathbf{x}^{(i)}$  as input, and a sequence  $\mathbf{y}$  as a potential output, the questions arises:

- How does an RNN produce  $\mathbf{y}_{(t)}$ , the  $t^{th}$  output?


Some choices

- Predict  $\mathbf{y}_{(t)}$  as a direct function of the prefix of  $\mathbf{x}$  of length  $t$ :

$$p(\mathbf{y}_{(t)} | \mathbf{x}_{(1)} \dots \mathbf{x}_{(t)})$$

**Direct function**



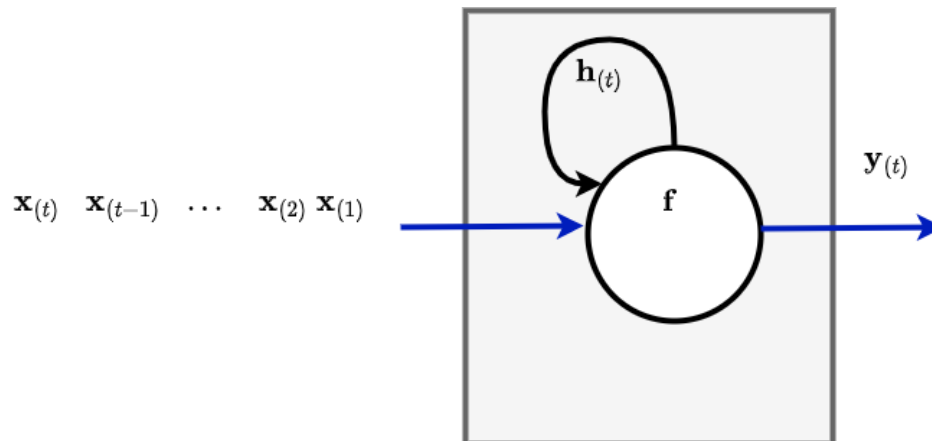


- Loop

- Uses a "latent state" that is updated with each element of the sequence, then predict the output

$$p(\mathbf{h}_{(t)} | \mathbf{x}_{(t)}, \mathbf{h}_{(t-1)}) \quad \text{latent variable } \mathbf{h}_{(t)} \text{ encodes } [\mathbf{x}_{(1)} \dots \mathbf{x}_{(t)}]$$
$$p(\mathbf{y}_{(t)} | \mathbf{h}_{(t)}) \quad \text{prediction contingent on latent variable}$$

Loop with latent state



## Latent state

The *latent state*  $\mathbf{h}_{(t)}$  is a kind of memory that acts as a *summary* of the prefix of sequence  $\mathbf{x}$  through time step  $t$ :

$$\mathbf{h}_{(t)} = \text{summary}(\mathbf{x}_{([1:t])})$$

Note that  $\mathbf{h}_{(t)}$  is a *vector* of fixed length.

Thus, it is a *fixed length* representation of the key aspects of a sequence  $\mathbf{x}$  of potentially *unbounded* length.

## Example

Let's use an RNN to compute the sum of a sequence numbers

- the latent state  $\mathbf{h}_{(t)}$  can be maintained as

$$\mathbf{h}_{(t)} = \text{summary}(\mathbf{x}_{([1:t])}) = \sum_{t'=1}^t \mathbf{x}_{(t')}$$

- by updating  $\mathbf{h}_{(t)}$  in the loop

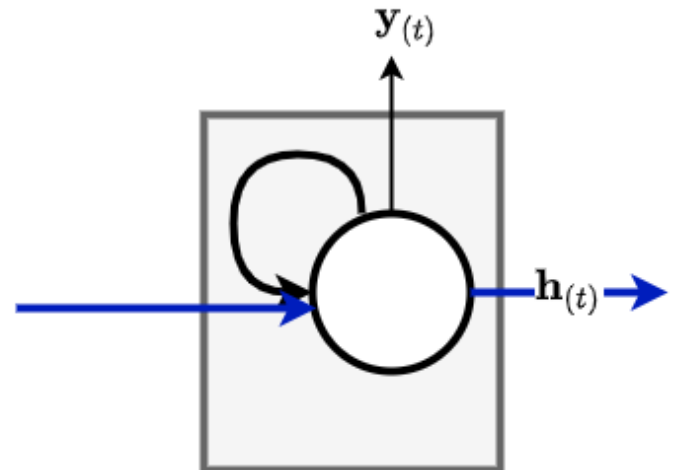
$$\mathbf{h}_{(t)} = \mathbf{h}_{(t-1)} + \mathbf{x}_{(t)}$$

Let's make this concrete with an example: a sequence of words

RNN

---

Machine Learning is easy not hard



$\mathbf{h}_{(t)}$  is a **fixed length** vector that "summarizes" the prefix of sequence  $\mathbf{x}$  up to element  $t$ .

The sequence is processed element by element, so order matters.

$$\mathbf{h}_{(0)} = \text{summary}([\text{Machine}])$$

$$\mathbf{h}_{(1)} = \text{summary}([\text{Machine}, \text{Learning}])$$

$$\vdots$$

$$\mathbf{h}_{(t)} = \text{summary}([\mathbf{x}_{(0)}, \dots, \mathbf{x}_{(t)}])$$

$$\vdots$$

$$\mathbf{h}_{(5)} = \text{summary}([\text{Machine}, \text{Learning}, \text{is}, \text{easy}, \text{not}, \text{hard}])$$

The importance of  $\mathbf{h}_{(t)}$  being *fixed length*

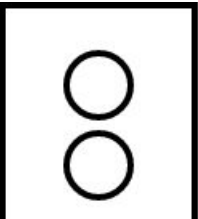
- can be used as input to other types of Neural Network layers
- which *don't* process sequences.

A typical example is a model for text classification (sentiment)

- Using an RNN to create a fixed length encoding of a variable length sequence
- A Head Layer that is a Binary Classifier

---

RNN Many to one; followed by classifier

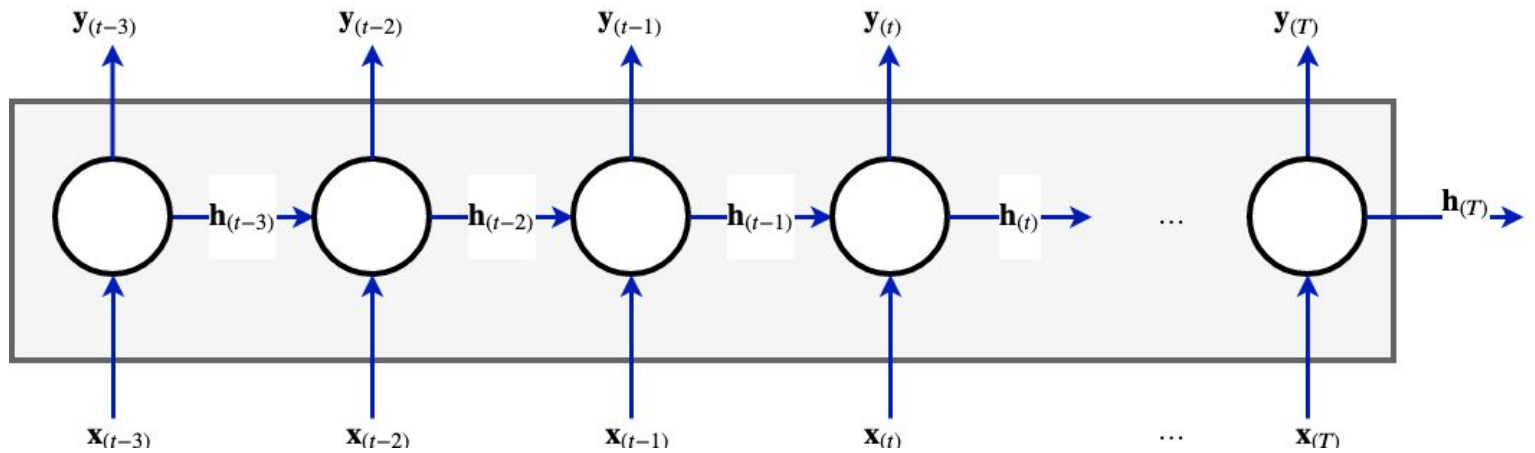


$\mathbf{y}_{(t)}$



# Unrolled RNN diagram

RNN many to many API



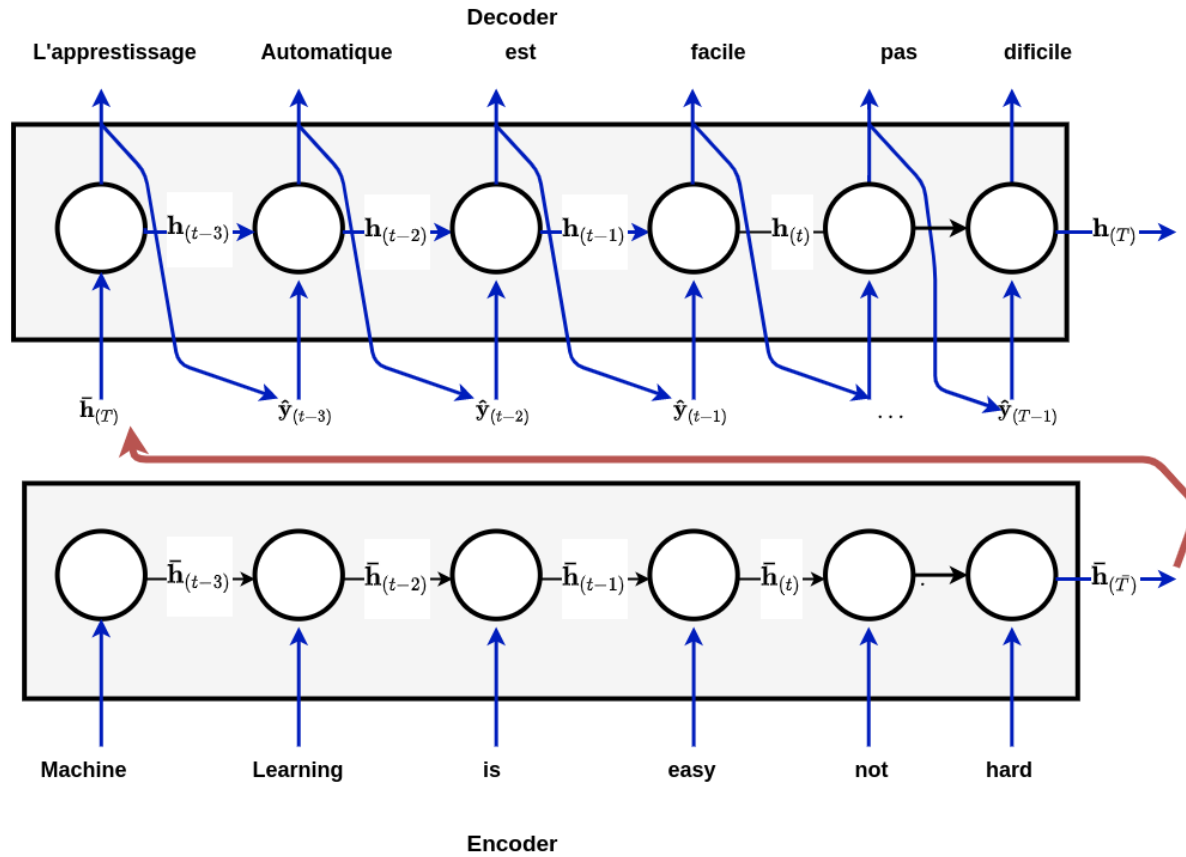
# Encoder-Decoder architecture; Auto-regressive

A very common architecture pairs two RNN's

- an Encoder, which summarizes the input sequence  $\mathbf{x}_{([1:\bar{T}])}$  via final latent state  $\bar{\mathbf{h}}_{(\bar{T})}$
- a Decoder, which takes the input summary  $\bar{\mathbf{h}}_{(\bar{T})}$  and outputs sequence  $\hat{\mathbf{y}}_{([1:T])}$

It is used for *Sequence to Sequence* tasks where both the input and output are sequences.

## Encoder-Decoder for language translation

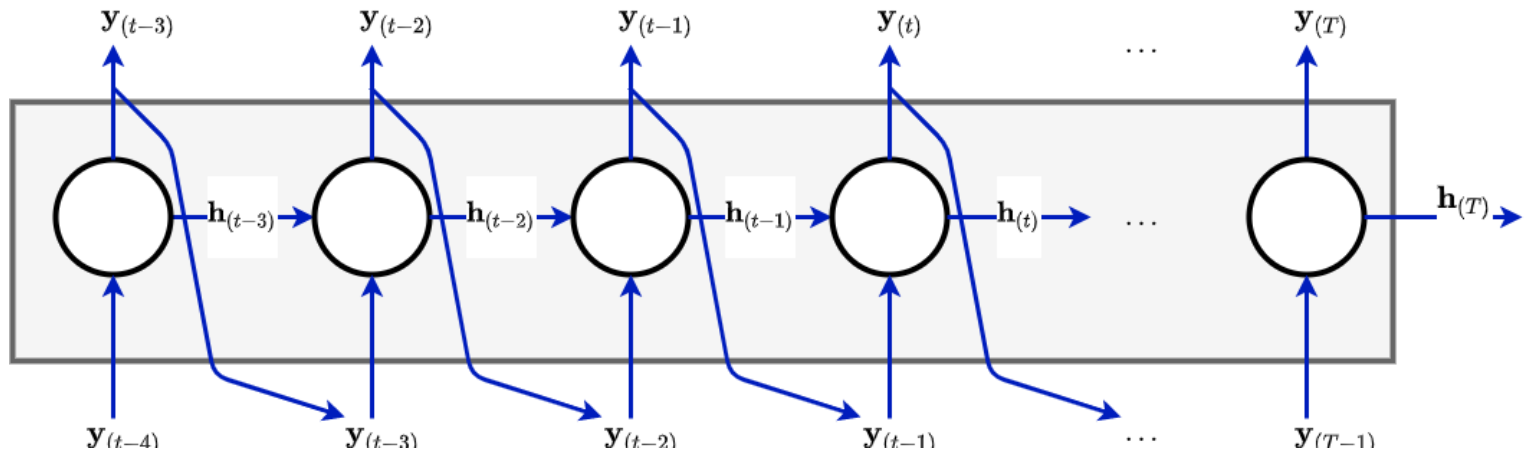


Notice that

- the output  $\hat{\mathbf{y}}_{(t-1)}$  of the Decoder at position  $(t - 1)$
- is used as the *input* at position  $t$

This is called *auto-regressive* behavior.

Test time: no forcing



In [2]: `print("Done")`

Done

