

Notation

To ensure that everyone is up to speed on notation, let's review

- [the notation \(ML Notation.ipynb\)](#) that we used in the "Classical Machine Learning" part of the intro course.
- [additional notation \(Intro to Neural Networks.ipynb\)](#) used in the "Deep Learning" part of the intro course

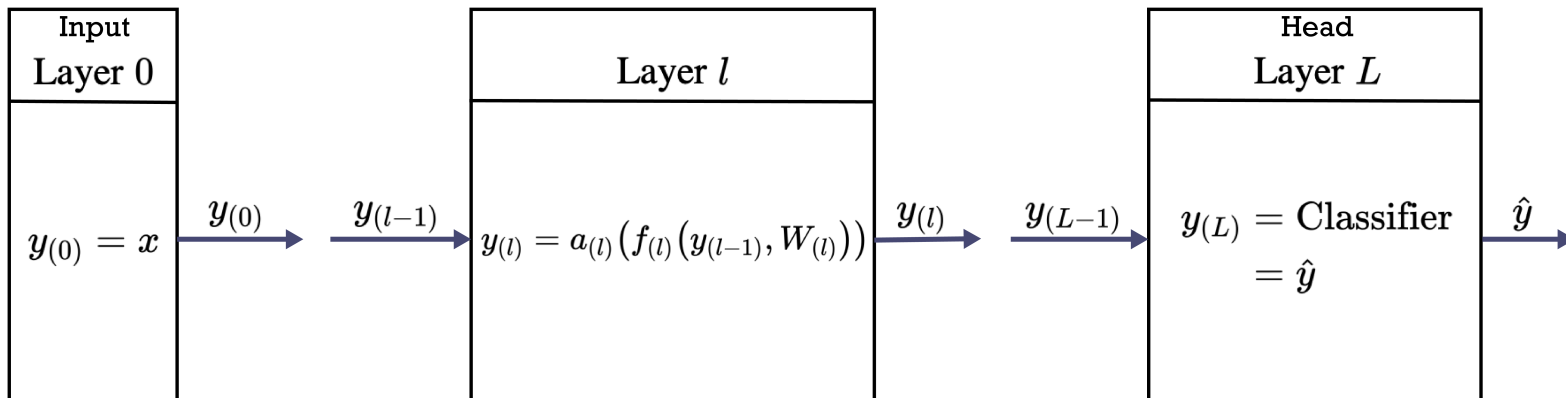
Representations

A path through a Neural Network can be viewed as a sequence of representation transformations

- transforming *raw features* $\mathbf{y}_{(0)} = \mathbf{x}$
- into *synthetic features* $\mathbf{y}_{(l)}$
 - varying with layer $1 \leq l \leq (L - 1)$
- of increasing abstraction

Thus, the output anywhere along the path is an *alternate representation* of the input

Path through a Neural Network



Shallow features are less abstract: "syntax", "surface"

Deeper features are more abstract: "semantics", "concepts"

- We may even interpret the features as "pattern matching" regions or concepts in the raw feature space.

For example, in a CNN

- shallow features are primitive shapes
- deeper features seem to recognize combinations of shallower features

Input features detected by layer

Layer 1



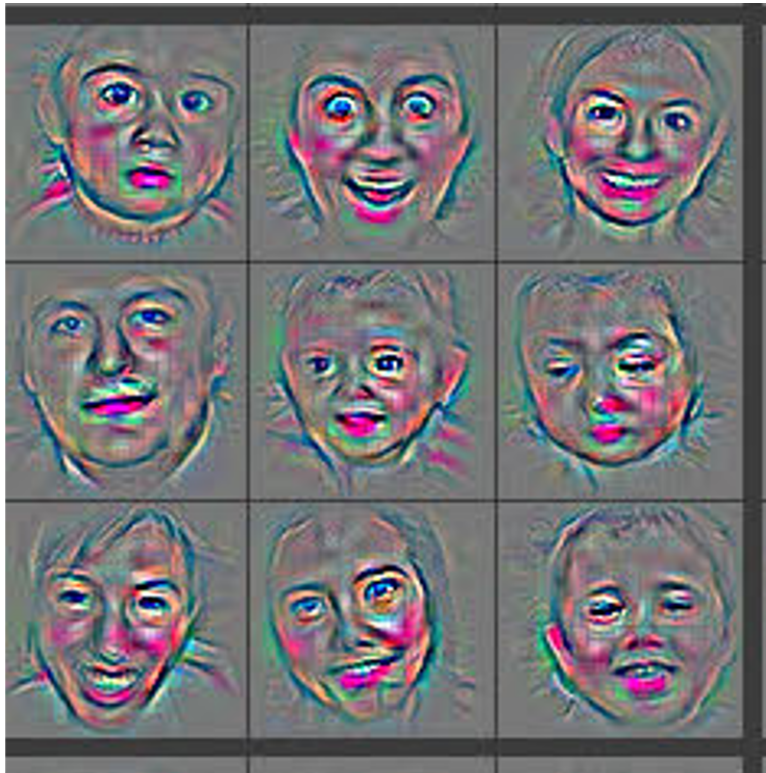
Layer 2



Layer 3



Saliency Maps and Corresponding Patches
Single Layer 5 Feature Map
On 9 Maximally Activating Input images



Layer 5 Feature Map (Row 11, col 1).

Attribution: <https://arxiv.org/abs/1311.2901> (<https://arxiv.org/abs/1311.2901>)

In the simple architectures of the Intro course, we mostly ignored the intermediate representations

$$\mathbf{y}_{(l)} : 1 \leq l \leq (L - 1)$$

The layers were referred to as "hidden" for a reason !

We will discover uses for intermediate representations and show how to build a "feature extractor" to obtain them from a given architecture.

In [2]: `print("Done")`

Done

