

# The Transformer: Code

We will examine a notebook that builds a miniature version of GPT: [tutorial view](https://keras.io/examples/generative/text_generation_with_miniature_gpt/)  
([https://keras.io/examples/generative/text\\_generation\\_with\\_miniature\\_gpt/](https://keras.io/examples/generative/text_generation_with_miniature_gpt/)).

- [Colab notebook](https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text_generation_with_miniature_gpt.ipynb) ([https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text\\_generation\\_with\\_miniature\\_gpt.ipynb](https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text_generation_with_miniature_gpt.ipynb))

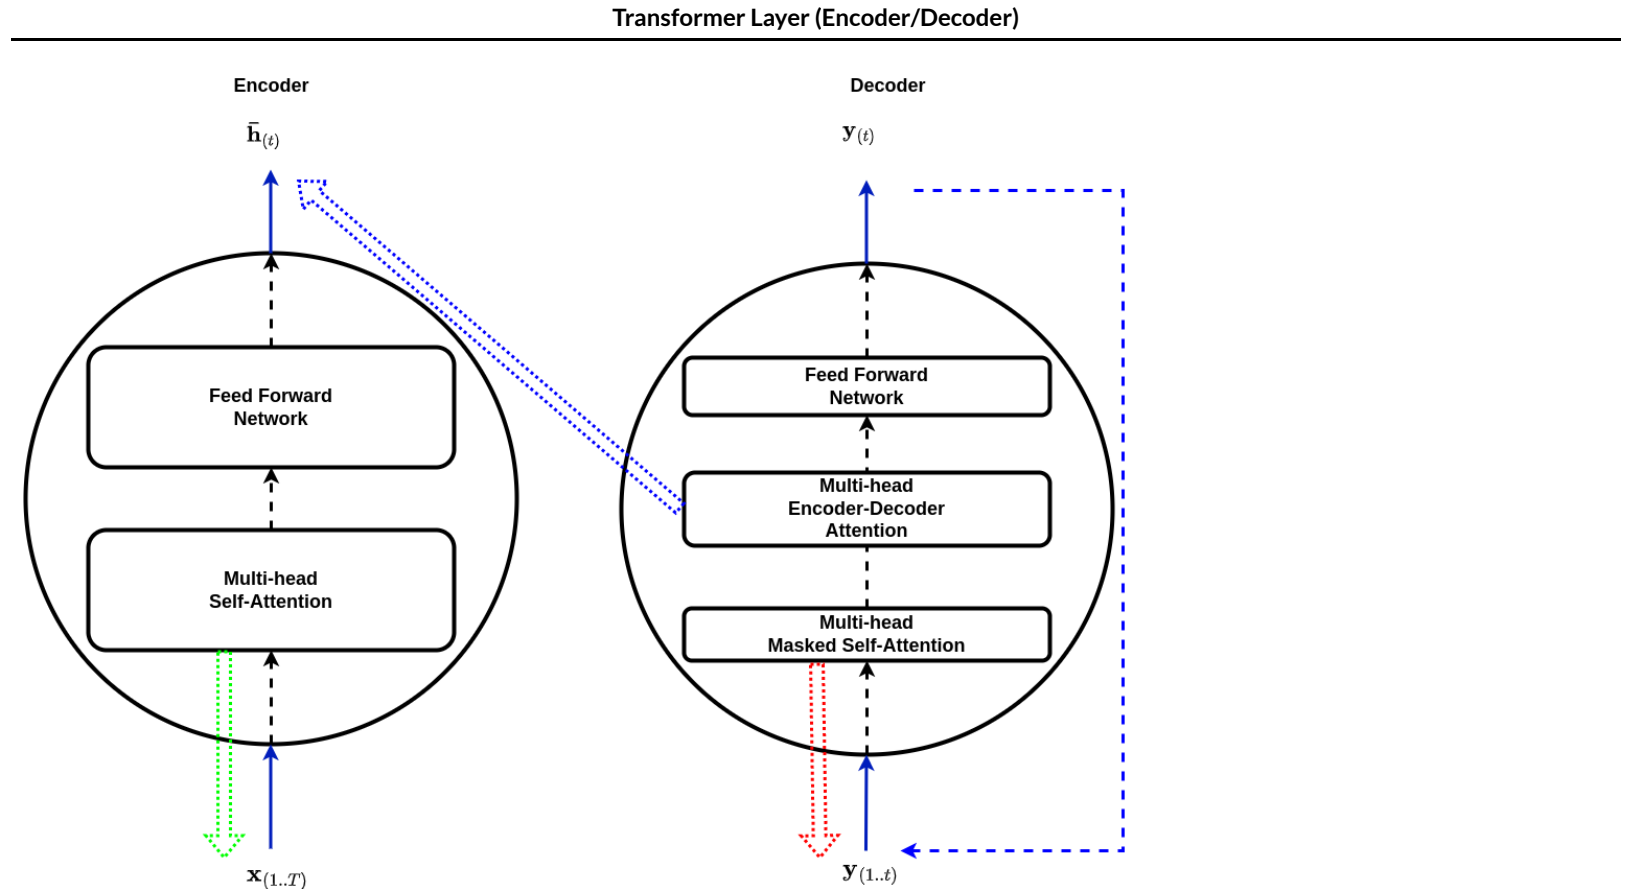
For an excellent tutorial on all the concepts, along with code, [see](https://www.tensorflow.org/text/tutorials/transformer)  
(<https://www.tensorflow.org/text/tutorials/transformer>).

---

## GPT-3 is a Decoder style Transformer

- autoregressive

Recall from our introduction to the Transformer (Encoder-Decoder)



The Decoder is the RHS of the image.

Here ([https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text\\_generation\\_with\\_miniature\\_gpt.ipynb#scroll-to=10](https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text_generation_with_miniature_gpt.ipynb#scroll-to=10)) we can see the Decoder

---

We first see a definition of the constants:

```
vocab_size = 20000 # Only consider the top 20k words
maxlen = 80 # Max sequence size
embed_dim = 256 # Embedding size for each token
num_heads = 2 # Number of attention heads
feed_forward_dim = 256 # Hidden layer size in feed forward network inside transformer
```

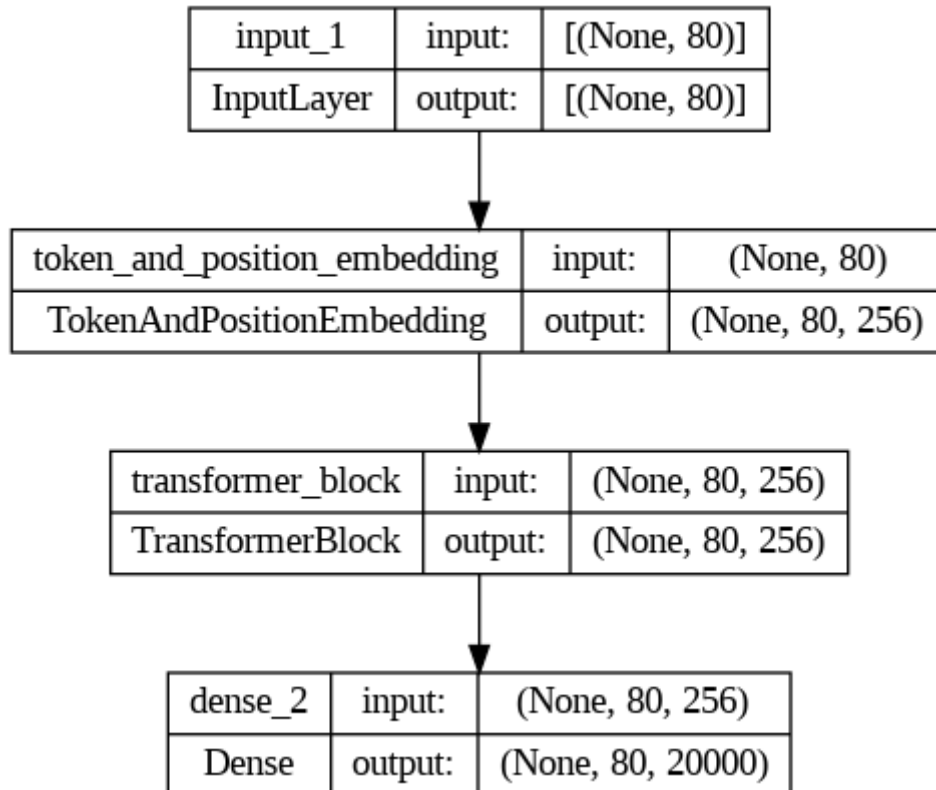
Relating the variable names to our notation

Notation	variable	value
$d_{\text{model}}$	embed_dim	256
$T$	max_len	80
$n_{\text{heads}}$	num_heads	2
	vocab_size	20,000

And the Decoder model:

```
def create_model():
    inputs = layers.Input(shape=(maxlen,), dtype=tf.int32)
    embedding_layer = TokenAndPositionEmbedding(maxlen, vocab_size, embed_dim)
    x = embedding_layer(inputs)
    transformer_block = TransformerBlock(embed_dim, num_heads, feed_forward_dim)
    x = transformer_block(x)
    outputs = layers.Dense(vocab_size)(x)
    model = keras.Model(inputs=inputs, outputs=[outputs, x])
    loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
    model.compile(
        "adam", loss=[loss_fn, None],
    ) # No loss and optimization based on word embeddings from transformer block
    return model
```

Here is the plot:



## Examining each layer

- Input
  - sequence (length  $T = 80$ ) of integers (index of a character within vocabulary)  $\mathbf{y}_{(1:T)}$
- TokenAndPositionEmbedding
  - maps sequence (length  $T = 80$ ) of integers (index of character)
  - into sequence (length  $T = 80$ ) of  $d_{\text{model}} = 256$  size representations
- TransformerBlock
  - maps sequence (length  $T = 80$ ) into sequence of latents  $\mathbf{h}_{(1:T)}$ 
    - one latent per position in input

- Dense
  - Classifier layer
  - maps sequence of latents
  - to sequence of probability vectors
    - each position is a probability vector of length `vocab_size`  
= 20000
    - position  $i$ : probability that output is element  $i$  of vocabulary
    - sum across positions in each vector is 100%



# Loss function

The `create_model` method also defines the Loss Function

```
loss_fn = tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True)
```

as Cross Entropy, as is common for a Classifier

Notice that the `SparseCategoricalCrossentropy` takes a vector (of length `vocab_size`) of **logits** rather than **probabilities**.

# TransformerBlock

Let's examine the [TransformerBlock](https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text_generation_with_miniature_gpt.ipynb#scroll=0) ([https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text\\_generation\\_with\\_miniature\\_gpt.ipynb#scroll=0](https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text_generation_with_miniature_gpt.ipynb#scroll=0)) in more detail

---

In [ ]:

```
class TransformerBlock(layers.Layer):
    def __init__(self, embed_dim, num_heads, ff_dim, rate=0.1):
        super().__init__()
        self.att = layers.MultiHeadAttention(num_heads, embed_dim)
        self.ffn = keras.Sequential(
            [layers.Dense(ff_dim, activation="relu"), layers.Dense(embed_dim)],
        )
        self.layernorm1 = layers.LayerNormalization(epsilon=1e-6)
        self.layernorm2 = layers.LayerNormalization(epsilon=1e-6)
        self.dropout1 = layers.Dropout(rate)
        self.dropout2 = layers.Dropout(rate)

    def call(self, inputs):
        input_shape = tf.shape(inputs)
        batch_size = input_shape[0]
        seq_len = input_shape[1]
        causal_mask = causal_attention_mask(batch_size, seq_len, seq_len, tf.bool)

        attention_output = self.att(inputs, inputs, attention_mask=causal_mask)

        attention_output = self.dropout1(attention_output)
        out1 = self.layernorm1(inputs + attention_output)
        ffn_output = self.ffn(out1)
        ffn_output = self.dropout2(ffn_output)
        return self.layernorm2(out1 + ffn_output)
```

We can see that the TransformerBlock is implemented as a Layer ( `layers.Layer` )

- so it will translate its input into output via a `call` method

The class `__init__` method defines the components of the Transformer

- stores them in instance variables:
  - Attention: `self.att`
  - Feed Forward Network FFN: `self.ffn`
  - Other: Layer Norms, Dropouts

The `call` method does the actual work

- Masked self-attention to  $\mathbf{y}_{(1:T)}$ 
  - Creates casual mask `causal_mask` to prevent peeking ahead at not-yet-generated output
    - `seq_len` is current length  $t$  of  $\mathbf{y}_{1:t}$
  - Attention block `self.att` applied to causally-masked input

```
attention_output = self.att(inputs, inputs,
                             attention_mask=causal_mask)
```
- Dropout `self.dropout1` and LayerNorm `layernorm1` applied to attention output
- Result passed through Feed Forward Network `self.ffn`

# TokenAndPositionEmbedding

Let's examine the [TokenAndPositionEmbedding](https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text_generation_with_miniature_gpt.ipynb#scroll=100) ([https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text\\_generation\\_with\\_miniature\\_gpt.ipynb#scroll=100](https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text_generation_with_miniature_gpt.ipynb#scroll=100)).

```
class TokenAndPositionEmbedding(layers.Layer):
    def __init__(self, maxlen, vocab_size, embed_dim):
        super().__init__()
        self.token_emb = layers.Embedding(input_dim=vocab_size, output_dim=embed_dim)
        self.pos_emb = layers.Embedding(input_dim=maxlen, output_dim=embed_dim)

    def call(self, x):
        maxlen = tf.shape(x)[-1]
        positions = tf.range(start=0, limit=maxlen, delta=1)
        positions = self.pos_emb(positions)
        x = self.token_emb(x)
        return x + positions
```

---

We can see that it too is implemented as a Layer.

The `call` method

- translates the input sequence
  - each position in the sequence is an integer index within the vocabulary
- into a sequence of pairs

- first element: token embedding

```
x = self.token_emb(x)
```

- second element: position embedding

```
positions = tf.range(start=0, limit=maxlen, delta=1)  
positions = self.pos_emb(positions)
```

As explained [in a prior module \(Transformer\\_PositionalEmbedding.ipynb#Representing-the-combined-token-and-positional-encoding\)](#).

- The output is not actually a sequence of *pairs*
  - it is a sequence of numbers
  - the token and positional emeddings are *added* not concatenated
    - concatenation would double the length
    - all layers in Transformer preserve output length equal input  
length =  $d_{\text{model}}$
- See the module's explanation as to why addition works



## Dense (Feed Forward Network)

We can see that the Feed Forward Network are two Dense layers

```
self.ffn = keras.Sequential(  
    [layers.Dense(ff_dim, activation="relu"), layers.Dense(embed_dim),]  
)
```

We may have been expecting the final layer of `TransformerBlock` to be outputting a probability vector (over the Vocabulary)

- a vector of length `vocab_size`
  - position  $i$  is probability that output is element  $i$  of the Vocabulary
- using a `softmax` activation
  - to make sure sum (across the `vocab_size` elements of the vector) of probabilities is 100%

But we see that the output is

- a singleton (not a vector)
- of size equal to `embed_dim =  $d_{\text{model}}$`

That is:

- the `Dense` component of the `TransformerBlock` is outputting the embedding of  $\hat{\mathbf{y}}_{(t)}$  rather than a probability vector

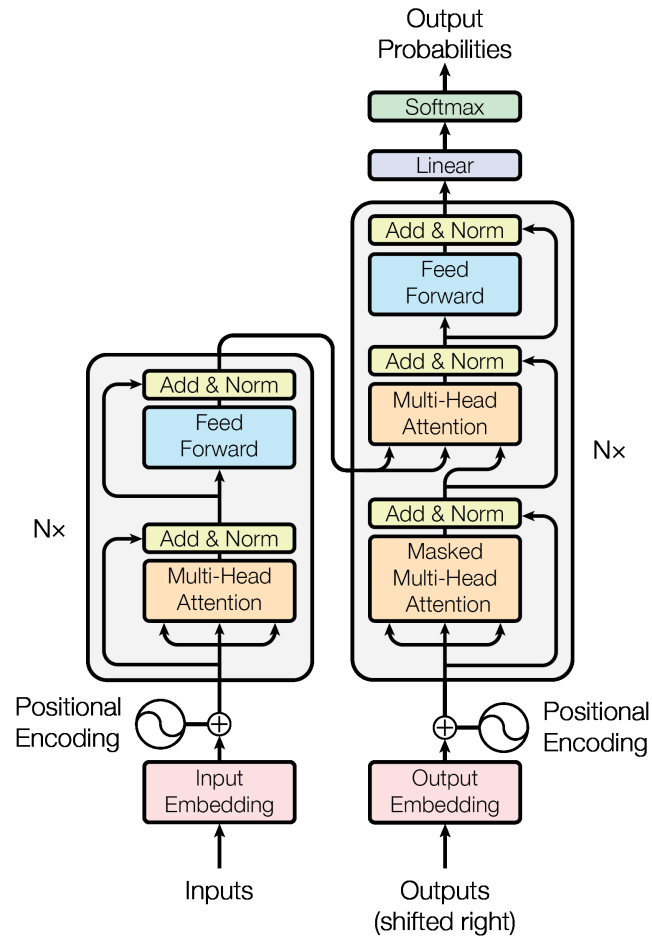
As we will see

- there is a layer in the Model *after* the TransformerBlock
- that produces the probability vector

# Skip connections

Here is a more detailed view of the Transformer

Transformer (Encoder/Decoder)



In particular, please focus on the arrows *into the "Add & Norm" layers*.

These are *skip connections* that bypass the Attention layers.

- *Residual Networks*

Where is this reflected in the code ?

It is a little subtle and easy to miss.

With the `call` method of the `TransformerBlock` please notice the statement

```
out1 = self.layer_norm1(inputs + attention_output)
```

- `inputs` is the input to the Attention layer

```
attention_output = self.att(inputs, inputs, attention_mask=causal_mask)
```

So the addition

`inputs + attention_output`

is joining (via addition)

- the output of the Attention layer
- the input of the Attention layer

This is the skip connection !

Similar code appears

```
ffn_output = self.ffn(out1)
ffn_output = self.dropout2(ffn_output)
return self.layer_norm2(out1 + ffn_output)
```

where

- the input to the FFN (i.e., `out1`)
- is joined (via addition) to the output of the FFN (i.e., `ffn_output`)

`out1 + ffn_output`



# Model

By examining the `create_model` function, we see that the output of the `TransformerBlock`

- is fed into a `Dense` layer
- which outputs a vector of length `vocab_size` (the correct length of a probability vector)
- and the output of this `Dense` layer is the output of the **model**
  - not the output of the `TransformerBlock`

```
outputs = layers.Dense(vocab_size)(x)
model = keras.Model(inputs=inputs, outputs=[outputs, x])
```
- Technically: the output vector is of *un-normalized logits* rather than probabilities

**- the logit vector can be turned into a probability vector via a `softmax`**

Thus, the Model outputs a vector of logits.

We can see how a token is sampled

- by converting the logit vector into a probability vector
- with the `sample_from` method of the `TextGenerator` callback

```
def sample_from(self, logits):
```

```
    logits, indices = tf.math.top_k(logits, k=self.k, sorted=True)
    indices = np.asarray(indices).astype("int32")
    preds = keras.activations.softmax(tf.expand_dims(logits, 0))[0]
    preds = np.asarray(preds).astype("float32")
    return np.random.choice(indices, p=preds)
```

Rather than outputting a probability vector

- which would require the user choosing one element from the vector (a word in the vocabulary)
- what is output is the *embedding* of the chosen word in the vocabulary

Since this output is compared against the correct label (i.e,  $\mathbf{y}_{(t+1)}$  for position  $t$ )

- we should also see that the *labels* used are embeddings

# Training

A `TextGenerator` ([https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text\\_generation\\_with\\_miniature\\_gpt.ipynb#scroll=0](https://colab.research.google.com/github/keras-team/keras-io/blob/master/examples/generative/ipynb/text_generation_with_miniature_gpt.ipynb#scroll=0)) call-back is used during training

- at the end every `self.print_every` epochs
  - a sample of  $\hat{\mathbf{y}}_{(1:T)}$  will be drawn
  - to illustrate what the model output would be up to that point in training
-

## The heart of the call-back

```
while num_tokens_generated <= self.max_tokens:
    ...
    y, _ = self.model.predict(x)
    sample_token = self.sample_from(y[0][sample_index])
    ...
```

- is a loop over positions  $t$
- that extends a fixed input (prefix of text) `start_tokens`
- to full length  $T$
- by sampling a token from the output for position  $t$

This is useful

- to see whether our model is learning as epochs advance
- to confirm the shape and type of the model output is a vector of logits
  - the model output for position  $t$ : `y, _ = self.model.predict(x)`
  - is passed to `sample_from`
  - which samples from the probability distribution derived from the logits (model output)

```
In [2]: print("Done")
```

Done

