

From PCA to VQ-VAE

[paper: vanilla VQ-VAE \(https://arxiv.org/pdf/1711.00937.pdf\)](https://arxiv.org/pdf/1711.00937.pdf)

[paper: VQ-VAE-2 \(https://arxiv.org/pdf/1906.00446.pdf\)](https://arxiv.org/pdf/1906.00446.pdf)

The common element in the design of any Autoencoder method is

- to create a latent representation \mathbf{z} of input \mathbf{x}
- such that \mathbf{z} can be (approximately) inverted to reconstruct \mathbf{x} .

Principal Components Analysis is a type of Autoencoder that produces a latent representation \mathbf{z} of \mathbf{x}

- \mathbf{x} is a vector of length n : $\mathbf{x} \in \mathbb{R}^n$
- \mathbf{z} is a vector of length $n' \leq n$: $\mathbf{z} \in \mathbb{R}^{n'}$

Usually $n' \ll n$: achieving *dimensionality reduction*

This is accomplished by decomposing \mathbf{x} into a weighted product of n *Principal Components*

- $\mathbf{V} \in \mathbb{R}^{n \times n}$

$$\mathbf{x} = \mathbf{z}' \mathbf{V}^T$$

- where $\mathbf{z}' \in \mathbb{R}^n$
- rows of \mathbf{V}^T are the components

So \mathbf{x} can be decomposed into the weighted sum (with \mathbf{z}' specifying the weights)

- of n component vectors
- each of length n

Since $\mathbf{z}' \in \mathbb{R}^n$: there is **no** dimensionality reduction just yet.

One can view \mathbf{V}^T as a kind of *code book*

- any \mathbf{x} can be represented (as a linear combination) of the *codes* (components) in \mathbf{V}^T

$$\mathbf{x} = \mathbf{z}' \mathbf{V}^T$$

\mathbf{z}' is like a translation of \mathbf{x} , using \mathbf{V} as the vocabulary.

- weights in the codebook
- rather than weights in the standard basis space $I \in \mathbb{R}^{n \times n} = \text{diagonal}(n)$

$$\mathbf{x} = \mathbf{x} I$$

Dimensionality reduction is achieved by defining \mathbf{z} as a length n' prefix of \mathbf{z}

- $\mathbf{z} = \mathbf{z}'_{1:n'}$
- $\mathbf{z} \in \mathbb{R}^{n'}$

Similarly, we needed only n' components from \mathbf{V}

- $\mathbb{V}^T = \mathbf{V}_{1:n'}^T$
- $\mathbb{V}^T \in \mathbb{R}^{n' \times n}$

We can construct an *approximation* $\hat{\mathbf{x}}$ of \mathbf{x} using *reduced dimension* \mathbf{z}' and \mathbb{V}

$$\hat{\mathbf{x}} = \mathbf{z}\mathbb{V}^T$$

The Autoencoder (and variants such as VAE) produces $\mathbf{z}^{(i)}$, the latent representation of $\mathbf{x}^{(i)}$

- directly
- independent of any other training example $\mathbf{x}^{(i')}$ for $i \neq i'$

Our goal in using AE's is in generating synthetic data

- the dimensionality reduction achieved thus far was a necessity, not a goal

Vector Quantized Autoencoder

A *Vector Quantized VAE* is a VAE with similarities to PCA. It creates \mathbf{z}

- which is an **integer**
- that is the index of a row
- in a codebook with K rows

That is: the input is represented by one of K possible vectors.

The goal is **not necessarily** dimensionality reduction.

Rather, there are some advantages to a **discrete** representation of a continuously-valued vector.

- Each vector
- Drawn from the infinite space of continuously-valued vectors of length n
- Can be approximated by one of K possible vectors of length n

Thus, a sequence of T continuously valued vectors

- can be represented as a sequence of T integers
- over a "vocabulary" defined by the code book

This is analogous to text

- sequence of words
- represented as a sequence of integer indices in a vocabulary of tokens

Once we put complex objects

- like images
- timeseries
- speech

into a representation similar to text

- we can have *mixed type* sequences
 - e.g., words, images

In a subsequent module we will take advantage of mixed type sequences

- to produce an image
- from a text *description* of the image
- using the "predict the next" element of a sequence technique of Large Language Models

DALL-E: Text to Image

Text input: "An illustration of a baby daikon radish in a tutu walking a dog"

Image output:



Details

Here is diagram of a VQ-VAE

- that creates a latent representation of a 3-dimensional image ($w \times h \times 3$)
- as a 2-dimensional matrix of integers

There is a bit of notation: referring to the diagram should facilitate understanding the notation.

VQ-VAE



In general, we assume the input has $\#S$ *spatial* dimensions

- where each location in the spatial dimension is a vector of length n
- input shape $(n_1 \times n_2 \dots \times n_{\#S} \times n)$

We will explain this diagram in steps.

First, we summarize the notation in a single spot for easy subsequent reference.

Notation summary

term	shape	meaning	
S	$(n_1 \times n_2 \dots \times n_{\#S})$	Spatial dimensions of $\#S$ -dimensional input	
\mathbf{x}	$\mathbb{R}^{S \times n}$	Input	
D		length of latent vectors (Encoder output, Quantized Encoder output, Codebook entry)	
\mathcal{E}		Encoder function	
$\mathbf{z}_e(\mathbf{x})$	$\mathbb{R}^{S \times D}$	Encoder output over each location of spatial dimension	
		$\mathbf{z}_e(\mathbf{x}) = \mathcal{E}(\mathbf{x})$	
$\mathbf{z}_e(\mathbf{x})$	\mathbb{R}^D	Encoder output at a single representative spatial location	
		$\mathbf{z}_e(\mathbf{x}) = \mathcal{E}(\mathbf{x})$	
K		number of codes	
\mathbf{E}	$\mathbb{R}^{K \times D}$	Codebook/Embedding	
		K codes, each of length D	
$e \in \mathbf{E}$	\mathbb{R}^D	code/embedding	
\mathbf{z}	$\{1, \dots, K\}^{S \times D}$	latent representation over all spatial dimensions	
\mathbf{z}	$\{1, \dots, K\}$	Latent representation at a single representative spatial location	
		one integer per spatial location	
$\lfloor \mathbf{z} \rfloor$		integer $\in [1 \dots K]$	Index k of $e_k \in \mathbf{E}$ that is closest to $\mathbf{z}_e(\mathbf{x})$
		$k = \underset{j \in [1, K]}{\operatorname{argmin}} \mathbf{z}_e(\mathbf{x}) - \mathbf{e}_j _2$	
		actually: encoded as a OHE vector of length K	
$\mathbf{z}_q(\mathbf{x})$	\mathbb{R}^D	Quantized $\mathbf{z}_e(\mathbf{x})$	
		$\mathbf{z}_q(\mathbf{x}) = e_k$ where $k = \lfloor \mathbf{z} \rfloor$	$\lfloor \mathbf{x} \rfloor$
		i.e, the element of codebook that is closest to $\mathbf{z}_e(\mathbf{x})$	
		$\mathbf{z}_q(\mathbf{x}) \approx \mathbf{z}_e(\mathbf{x})$	
$\tilde{\mathbf{x}}$	n	Output: reconstructed \mathbf{x}	

term	shape	meaning
		$\text{pr}\{ \mathbf{x}$ $\mathbf{z}_q(\mathbf{x}) \}$
\mathcal{D}	$\mathbb{R}^{n'} \rightarrow \mathbb{R}^n$	Decoder

Quantization

Let S denote the spatial dimensions, e.g. $S = (n_1 \times n_2)$ for 2D

So input $\mathbf{x} \in \mathbb{R}^{S \times n}$

- n features over S spatial locations

The input \mathbf{x} is transformed in a sequence of steps

- Encoder output (continuous value)
- Latent representation (discrete value)
 - Quantized (continuous value)

In the first step, the *Encoder* maps input \mathbf{x}

- to Encoder output $z_e(\mathbf{x})$
- an alternate representation of D features over S' spatial locations

(For simplicity, we will assume $S' = S$)

Notational simplification

In the sequel, we will apply the same transformation **to each element** of the spatial dimension

Rather than explicitly iterating over each location we write

$$\mathbf{z}_e(\mathbf{x}) \in \mathbb{R}^D$$

to denote a representative element of $\mathbf{z}_e(\mathbf{x})$ at a single location $s = (i_1, \dots, i_{\#S})$

$$\mathbf{z}_e(\mathbf{x}) = \mathbf{z}_e(\mathbf{x})_s$$

We will continue the transformation at the single representative location

- and implicitly iterate over all locations $s \in S$

The continuous (length D) Encoder output vector $\mathbf{z}_e(\mathbf{x})$

- is mapped to a *latent representation* $q(\mathbf{z}|\mathbf{x})$
- which is a **discrete** value (integer)

$$k = q(\mathbf{z}|\mathbf{x}) \in \{1, \dots, K\}$$

where k is the *index* of a row \mathbf{e}_k in codebook \mathbf{E}

$$\mathbf{e}_k = \mathbf{E}_k \in \mathbb{R}^D$$

The codebook is also called an *Embedding* table.

k is chosen such that \mathbf{e}_k is the row in \mathbf{E} closest to $\mathbf{z}_e(\mathbf{x})$

$$\begin{aligned} k &= q(\mathbf{z}|\mathbf{x}) \\ &= \operatorname{argmin}_{j \in \{1, \dots, K\}} \|\mathbf{z}_e(\mathbf{x}) - \mathbf{e}_j\|_2 \end{aligned}$$

We denote the codebook vector

- closest to representative encoder output $\mathbf{z}_e(\mathbf{x})$
 - as $\mathbf{z}_q(\mathbf{x})$
- $$\mathbf{z}_q(\mathbf{x}) = \mathbf{e}_k \text{ where } k = q(\mathbf{z}|\mathbf{x})$$

The Decoder tries to invert the codebook entry $\mathbf{e}_k = \mathbf{z}_q(\mathbf{x})$ so that

$$\begin{aligned}\tilde{\mathbf{x}} &= \mathcal{D}(\mathbf{z}_q(\mathbf{x})) \\ &\approx \mathbf{x}\end{aligned}$$

Discussion

Why do we need the CNN Encoder ?

The input \mathbf{x} is first transformed into an *alternate representation*

- the **number** and shape of the spatial dimensions are preserved (not necessary)
- but the number of features is transformed from n raw features to $D \geq n$ synthetic features
 - typical behavior for, e.g., an image classifier

The part of the VQ-VAE after the initial CNN

- reduces the size of the **feature dimension** from D to 1
- this is the primary source of dimensionality reduction
 - the raw n of image input is usually only $n = 3$ channels

It may be useful for the CNN to *down-sample* spatial dimension S to a smaller S'

For example

- 3 layers of stride 2 CNN layers
- will reduce a 2D image of spatial dimension $(n_1 \times n_2)$
- to spatial dimension $(\frac{n_1}{8} \times \frac{n_2}{8})$

This replaces each $(8 \times 8 \times n)$ *patch* of raw input

- into a single vector of length D
- that summarizes the (8×8) the patch

One possible role (not strictly necessary) for the CNN Encoder

- is to replace a large spatial dimensions
- by smaller "summaries" of local neighborhoods (patches)

Why quantize ?

Quantization

- converts the continuous $\mathbf{z}_e(\mathbf{x})$
- into discrete $q(\mathbf{z}|\mathbf{x})$
- representing the approximation $\mathbf{z}_q(\mathbf{x}) \approx \mathbf{z}_e(\mathbf{x})$

The Decoder inverts the approximation.

Why bother when the Quantization/De-Quantization is Lossy ?

One motivation comes from observing what happens if we *quantize and flatten* the $\#S'$ -dimensional spatial locations to a one-dimensional vector.

Quantizing replaces each patch with a single integer index.

- the integer is the index of an *image token* within a list of K possible tokens

By flattening the quantized higher dimensional matrix of patches, we convert the input

- into a sequence of image tokens
- over a "vocabulary" defined by the codebook \mathbf{E} .

This yields an image representation

- similar to the representation of text

Thus, we open the possibility of processing sequences of mixed text and image tokens.

Quantized image embeddings mixed with Text: preview of DALL-E

The Large Language Model operates on a sequence of text tokens

- where the text tokens are fragments of words
- when run autoregressively
 - concatenating each output to the initial input sequence
 - the LLM shows an ability to produce a "sensible" continuation of an initial "thought"

Suppose we train a LLM on input sequences

- that start with a sequence of *text* tokens describing an image
- followed by a separator [SEP] token
- followed by a sequence of of quantized image tokens

<text token> <text token> ... <text token> [SEP] <image token> <image token> ...

What continuation will our trained LLM produce given prompt

<text token> <text token> ... <text token> [SEP]

Hopefully:

- a sequence of *image tokens*
- that can be reconstructed
- into an image matching the description given by the text tokens !

That is the key idea behind a Text to Image model called DALL-E that we will discuss in a later module.

There remains an important technical detail

- the embedding space of text and image are distinct
- they need to be merged into a common embedding space

We will visit these issues in the module on CLIP.

Loss function

The Loss function for the VQ-VAE entails several parts

- Reconstruction loss
 - enforcing constraint that reconstructed image is similar to input
$$\tilde{\mathbf{x}} \approx \mathbf{x}$$
- Vector Quantization (VQ) Loss:
 - enforcing similarity of quantized encoder output and actual encoder output
$$\mathbf{z}_q(\mathbf{x}) \approx \mathbf{z}_e(\mathbf{x})$$
- Commitment Loss
 - a constraint that prevents the Quantization of $\mathbf{z}_e(\mathbf{x})$ from alternating rapidly between code book entries

Reconstruction Loss

The Reconstruction Loss term is our familiar: Maximize Likelihood

- written to minimize the negative of the log likelihood, as usual

$$p(\mathbf{x}|\mathbf{z}_q(\mathbf{x}))$$

The Decoder is responsible for the Reconstruction Loss (through the term $\tilde{\mathbf{x}}$)

- Straight Through Estimation (which we shall subsequently discuss)
- means the gradients from Decoder input $\mathbf{z}_q(\mathbf{x})$ flows directly to Encoder output $\mathbf{z}_e(\mathbf{x})$
- *without* affecting the Embeddings
 - i.e., the embedding $e_k \in \mathbb{E}$ (where $k = q(\mathbf{z}|\mathbf{x})$) is **not** affected by the Reconstruction Loss

Straight Through Estimation (discussed below) causes the gradient from Reconstruction Loss to "by-pass" \mathbb{E}

- effectively, for the purpose of gradient/weight update:

$$\mathbf{z}_q(\mathbf{x}) = \mathbf{z}_e(\mathbf{x})$$

The role of the other Loss terms (e.g., Vector Quantizer Loss) are to ensure that the Embeddings \mathbb{E} get updated.

Vector Quantization Loss

The Vector Quantization Loss and Commitment Loss are similar.

Vector Quantization Loss:

$$\| \text{sg}(\mathbf{z}_e(\mathbf{x})) - \mathbf{z}_q(\mathbf{x}) \|$$

where sg is the *Stop Gradient* Operator (details to follow).

The purpose of the Vector Quantization Loss is to "learn" the Embedding (codebook) \mathbf{E}

- by moving e_k closer to Encoder output $\mathbf{z}_e(\mathbf{x})$

Commitment Loss

Commitment Loss:

$$\|\mathbf{z}_e(\mathbf{x}) - \text{sg}(\mathbf{z}_q(\mathbf{x}))\|$$

It is similar to the Vector Quantization loss except for the placement of the Stop Gradient operator.

It's purpose is ensure that the embeddings \mathbb{E} converge.

Total Loss

Loss function

$$\begin{aligned}\mathcal{L}(\mathbf{x}, \mathcal{D}(\mathbf{e})) = & \|\mathbf{x} - \mathcal{D}(\mathbf{e})\|_2^2 && \text{Reconstruction Loss} \\ & + \|\text{sg}[\mathcal{E}(x)] - \mathbf{e}\|_2^2 && \text{VQ loss, codebook loss: train codebook} \\ & + \beta \|\text{sg}[\mathbf{e}] - \mathcal{E}(\mathbf{x})\|_2^2 && \text{Commitment Loss: force } E(\mathbf{x}) \text{ to be clo} \\ & \text{where } \mathbf{e} = \mathbf{z}_q(\mathbf{x})\end{aligned}$$

Need the stop gradient operator sg to control the mutual dependence

- of $\mathcal{E}(\mathbf{x})$ and \mathbf{e}
-

Stop Gradient operator

The `sg` operator is the *Stop Gradient* operator.

On the Forward Pass, it acts as an Identity operator

$$\text{sg}(\mathbf{x}) = \mathbf{x}$$

But on the Backward Pass of Backpropagation: *it stops the gradient* from flowing backwards

$$\frac{\partial \text{sg}(\mathbf{x})}{\partial \mathbf{y}} = 0 \text{ for all } \mathbf{y}$$

Why is this operator necessary for the VQ-VAE ?

The problem lies in the Quantization operation

$$\begin{aligned} k &= q(\mathbf{z}|\mathbf{x}) \\ &= \operatorname{argmin}_{j \in \{1, \dots, K\}} \|\mathbf{z}_e(\mathbf{x}) - \mathbf{e}_j\|_2 \end{aligned}$$

argmin is not differentiable !

- small changes in the arguments cause a change from k to k'
- It is not continuous at the point that its value switches between k and $k' \neq k$
- It is not necessarily deterministic:
 - when $\mathbf{e}_k = \mathbf{e}_{k'}$ for $k \neq k'$

The reason it appears in the Vector Quantization Loss and Commitment Loss is to prevent a feedback loop

- Encoder updating $\mathbf{z}_e(\mathbf{x})$ reduces Reconstruction Loss *assuming* embeddings remain constant
- But changing Encoder output results in embeddings being updated
- So embeddings *do not* remain constant
- The net effect may not be a reduction in Reconstruction Loss

The Stop Gradient of the Vector Quantization Loss prevents a change in the Encoder weights (and thus, $\mathbf{z}_e(\mathbf{x})$) from affecting the embeddings.

The Stop Gradient in the Commitment Loss prevents a change in the Embeddings from affecting the Encoder weights (and thus, $\mathbf{z}_e(\mathbf{x})$).

This prevents a feedback loop.

Straight through Estimation and the Stop Gradient operator `sg`

We can see the Stop Gradient operator in action by examining the code of the `VectorQuantizer` layer.

It is a tool that helps us implement a technique called *Straight Through Estimation*.

Let's recall the definition of the Loss Gradient used in Backpropagation

$$\mathcal{L}'_{(l)} = \frac{\partial \mathcal{L}}{\partial \mathbf{y}_{(l)}}$$

It is the derivative of \mathcal{L} with respect to the output of layer l , i.e., $\mathbf{y}_{(l)}$.

Although we state this with respect to a "layer-ed" architecture this is for notational convenience only

- the same is true if we replace "layer" with "operator" whose input is denoted $\mathbf{y}_{(l-1)}$ and output denoted $\mathbf{y}_{(l)}$

Back propagation inductively updates the Loss Gradient from the output of layer l to its inputs (e.g., prior layer's output $\mathbf{y}_{(l-1)}$)

- Given $\mathcal{L}'_{(l)}$
- Compute $\mathcal{L}'_{(l-1)}$
- Using the chain rule

$$\begin{aligned}\mathcal{L}'_{(l-1)} &= \frac{\partial \mathcal{L}}{\partial \mathbf{y}_{(l-1)}} \\ &= \frac{\partial \mathcal{L}}{\partial \mathbf{y}_{(l)}} \frac{\partial \mathbf{y}_{(l)}}{\partial \mathbf{y}_{(l-1)}} \\ &= \mathcal{L}'_{(l)} \frac{\partial \mathbf{y}_{(l)}}{\partial \mathbf{y}_{(l-1)}}\end{aligned}$$

The loss gradient "flows backward", from $\mathbf{y}_{(L+1)}$ to $\mathbf{y}_{(1)}$.

This is referred to as the *backward pass*.

That is:

- the upstream Loss Gradient $\mathcal{L}'_{(l)}$
- is modulated by the local gradient $\frac{\partial \mathbf{y}_{(l)}}{\partial \mathbf{y}_{(l-1)}}$
- where the "layer" is the operation transforming input $\mathbf{y}_{(l-1)}$ to output $\mathbf{y}_{(l)}$

What happens when the operation implemented by the function that takes $\mathbf{y}_{(l-1)}$ to $\mathbf{y}_{(l)}$ is either

- non-differentiable
- or has zero derivative almost everywhere
- non-deterministic (e.g., `tf.argmax` when two inputs are identical)

In the Quantization, we use `tf.argmax`.

If

- $\frac{\partial \mathbf{y}_{(l)}}{\partial \mathbf{y}_{(l-1)}} = 0$
- then the update

$$\mathcal{L}'_{(l-1)} = \mathcal{L}_{(l)} * 0 = 0$$

- the quantization operation disconnects the gradient flow from the Decoder backwards to the Encoder.
- Encoder won't learn

The solution is the Straight Through Estimator:

- identity operation on forward pass
- with local derivative $\frac{\partial \mathbf{y}_{(l)}}{\partial \mathbf{y}_{(l-1)}}$ **defined** to be equal to 1

We see this in the [Colab \(https://keras.io/examples/generative/vq_vae/\)](https://keras.io/examples/generative/vq_vae/) implementation of Vector Quantization (the VectorQuantizer layer)

```
class VectorQuantizer(layers.Layer):  
    ...  
    def call(self, x):  
    ...  
        # Straight-through estimator.  
        quantized = x + tf.stop_gradient(quantized - x)
```

Code similar to the [VectorQuantizer of the paper's authors \(https://github.com/deepmind/sonnet/blob/v1/sonnet/python/modules/nets/vqvae.py\)](https://github.com/deepmind/sonnet/blob/v1/sonnet/python/modules/nets/vqvae.py)

The last line is a "straight through estimator"

(<https://www.hassanaskary.com/python/pytorch/deep%20learning/2020/09/19/intuitive-explanation-of-straight-through-estimators.html>)

For this VectorQuantizer "layer" l

- given input x the forward pass returns quantized
$$\begin{array}{l} y_{(l-1)} = \text{quantized}(x) \\ \text{layer input} = x + \text{tf.stop_gradient}(\text{quantized} - x) \\ \text{layer output} = \text{quantized} \end{array}$$
 since $\text{tf.stop_gradient}(\text{quantized} - x)$ is zero

On the backward pass: the local gradient $\frac{\partial \mathbf{y}_{(l)}}{\partial \mathbf{y}_{(l-1)}}$ is

$$\begin{aligned} \frac{\partial \mathbf{y}_{(l)}}{\partial \mathbf{y}_{(l-1)}} &= \frac{\partial \mathbf{x} + \text{tf.stop_gradient}(\text{quantized} - \mathbf{x})}{\partial \mathbf{x}} && \text{since } \mathbf{y}_{(l)} = \mathbf{x} + \text{tf.stop_gradient}(\text{quantized} - \mathbf{x}) \\ &&& \text{and } \mathbf{y}_{(l-1)} = \mathbf{x} \\ &= \frac{\partial \mathbf{x}}{\partial \mathbf{x}} + \frac{\partial \text{tf.stop_gradient}(\text{quantized} - \mathbf{x})}{\partial \mathbf{x}} && \text{derivative of sum is sum of derivatives} \\ &= 1 && \text{since } \frac{\partial \text{tf.stop_gradient}(\text{..})}{\partial \mathbf{x}} = 0 \text{ by definition} \end{aligned}$$

Thus, for the VectorQuantizer "layer" l

$$\begin{aligned}\mathcal{L}'_{(l-1)} &= \frac{\partial \mathcal{L}}{\partial \mathbf{y}_{(l-1)}} \\ &= \mathcal{L}'_{(l)} \frac{\partial \mathbf{y}_{(l)}}{\partial \mathbf{y}_{(l-1)}} \\ &= \mathcal{L}'_{(l)}\end{aligned}$$

The Straight Through Estimator passes the Loss Gradient backwards unchanged

- because the `tf.stop_gradient` **kills** the gradient along one path

In the VQ-VAE, straight through estimation

- passes the gradient from the Decoder input back to the Encoder outputs
- ignoring the quantization
- allowing the Encoder to adapt to reduce Reconstruction Loss

Learning the distribution of latents

For a VAE, we assume a functional form for the prior distribution of latents $q(\mathbf{z})$

- a *convenient* choice is Normal. See [our VAE module](#)
([VAE_Generative.ipynb#Choosing- \$q\(\mathbf{z}\)\$](#)).

The authors wish to do away with an assumption of the prior distribution $q(\mathbf{z})$.

Retaining spatial/temporal dimensions in $\mathbf{z}_q(\mathbf{x})$ is key to achieving this goal.

The authors *flatten*

- the spatial/temporal dimensions $(n_1 \times n_2 \dots \times n_{\#S})$ of \mathbf{Z}
- into a sequence

$$\mathbf{Z}_{(1)}, \mathbf{Z}_{(2)}, \dots, \mathbf{Z}_{(n_1 * n_2 \dots * n_{\#S})}$$

For example: for two spatial dimensions $(h \times w)$

- $\mathbf{Z}_{(k)}$
- is the quantization of $\mathbf{Z}_c^{(r)}$
- the element at row r column c
- for $r = \text{int}(\frac{k}{w}), c$
 $= (k \bmod w)$

The authors then learn an autoregressive model for sequences

$$p(\mathbf{z}_{(k+1)} | \mathbf{z}_{(1)}, \dots, \mathbf{z}_{(k)})$$

by using some Autoregressive model (e.g, PixelCNN) to predict $\mathbf{z}_{(k+1)}$ from its predecessors in the sequence.

We are familiar with this Autoregressive model (in the case of NLP) as the Language Model objective.

Unlike the "convenient" common choice of Normality for the VAE

- the Autoregressive model does **not** assume the type of the distribution

The conditional distribution

$$p(\mathbf{z}_{(k+1)} | \mathbf{z}_{(1)}, \dots, \mathbf{z}_{(k)})$$

of each element

- is learned
- is easy to sample
 - seed the model with $\mathbf{z}_{(1)}$, generate the rest of the sequence one element at a time

In [2]: `print("Done")`

Done

