

# Wasserstein GAN: Motivation

[paper \(https://arxiv.org/pdf/1701.07875.pdf\)](https://arxiv.org/pdf/1701.07875.pdf)

To summarize what we have learned about standard GANs:

- Adversarial Training minimizes the Jensen Shannon Distance between  $p_{\text{model}}$  and  $p_{\text{data}}$
- They have the reputation for being difficult to train
  - A Discriminator that is too good, too soon inhibits the ability of the Generator to learn to generate
  - The Generator may "mode collapse" and not produce a variety of outputs

The *Wasserstein GAN (WGAN)* is a pair of Neural Networks NN:

- the Generator
- the Discriminator
  - technically, this is a "critic"
    - rather than producing a probability of "Real"
    - it produces a "score" measuring how real the input is
      - larger negative: more real
      - larger positive: less real

The pair is trained to minimize an *approximation* of

$$\mathbb{W}(p_{\text{data}}, p_{\text{model}})$$

where  $\mathbb{W}$  is the *Wasserstein Distance*, also known as the *Earth Move Distance (EMD)* measure.

# Earth Move Distance (EMD)

## Aside

You need some knowledge of Measure Theory to understand the math.

In the absence, there are two good blogs I recommend in order to get a flavor

- [Sorta Insightful \(https://www.alexirpan.com/2017/02/22/wasserstein-gan.html\)](https://www.alexirpan.com/2017/02/22/wasserstein-gan.html).
- [Wen \(https://arxiv.org/pdf/1904.08994.pdf\)](https://arxiv.org/pdf/1904.08994.pdf).

Like the KL and Jensen-Shannon Distances, the EMD is a measure of the difference between two distributions.  $p_{\text{data}}$  and  $p_{\text{model}}$ .

It has an intuitive explanation

*The minimum amount of "work" involved in moving probability mass between the two distributions in order to make them identical*

"Work" means: the product of

- the quantity  $\gamma(x, y)$  of the mass moved from  $x$  to  $y$
- and the distance  $\|x - y\|$  it is moved

We can easily illustrate with two discrete distributions (example from [Wen](https://arxiv.org/pdf/1904.08994.pdf) [.https://arxiv.org/pdf/1904.08994.pdf](https://arxiv.org/pdf/1904.08994.pdf)).

Let  $P, Q$  be the two distributions, represented as vectors since there are discrete and measured over the same indices.

$$\begin{aligned} P &= [3, 2, 1, 4] \\ Q &= [1, 2, 4, 3] \end{aligned}$$

$P_i$  (resp.,  $Q_i$ ) is the probability as  $i$  in each of the distributions, for  $1 \leq i \leq 4$ .

For illustration, we will move

- a quantity  $\delta_i$  of probability in  $P$  between adjacent indices  $(i - 1)$  and  $i$
- in order to make  $P_i = Q_i$  ( $Q_i$  remains fixed)
- The distance is 1 (and hence work is equal to quantity moved).

We can define  $\delta_i$  recursively:

$$\begin{aligned}\delta_0 &= 0 \\ \delta_{i+1} &= \delta_i + P_i - Q_i\end{aligned}$$

That is, the amount  $\delta_{i+1}$  moved from  $P_i$  in order to make  $P_i = Q_i$  is

- the difference  $(P_i - Q_i)$  between original value of  $P_i$  and  $Q_i$
- plus the additional quantity  $\delta_i$  that was moved into  $P_i$

$$\delta_0 = 0$$

$$\delta_1 = 0 + 3 - 1 = 2$$

$$\delta_2 = 2 + 2 - 2 = 2$$

$$\delta_3 = 2 + 1 - 4 = -1$$

$$\delta_4 = -1 + 4 - 3 = 0$$

Work is positive so taking absolute values

$$\mathbb{W}(P, Q) = \sum_{i=1}^4 1 * |\delta_i| = 5$$

For continuous distributions

$$\mathbb{W}(p_{\text{data}}, p_{\text{model}}) = \inf_{\gamma \in \Pi(p_{\text{data}}, p_{\text{model}})} \mathbb{E}_{(x,y) \sim \gamma} \|x - y\|$$

where

- $\Pi(p_{\text{data}}, p_{\text{model}})$  are the set of possible joint distributions with marginal  $p_{\text{data}}$  and  $p_{\text{model}}$
- $\gamma$  is a quantity to move from  $x$  to  $y$  (for all  $x, y$ )
  - distance between  $x$  and  $y$  is  $\|x - y\|$
- $\inf$  is the infimum (Greatest Lower Bound)



## Approximation of $\mathbb{W}(p_{\text{data}}, p_{\text{model}})$

**Warning:** the math is stated without much explanation

The infimum is intractable (or at least: not practical to compute).

Equation 2 in the paper states that for certain functions  $f$ , the distance is also equal to

$$\mathbb{W}(p_{\text{data}}, p_{\text{model}}) = \inf_{\|f\|_2 \leq 1} \mathbb{E}_{x \sim p_{\text{data}}} f(x) - \mathbb{E}_{x \sim p_{\text{model}}} f(x)$$

One can look at  $f$  as a "score" of  $x$  being "Real" (not fake) where

- a high negative score is a highly confident "Real"
- a high positive score is a highly confident "Fake"

The goal is

- for function  $f$  to create a *large spread* between scores of Real and Fake.
- for function  $f$  to be *approximated* by the Discriminator  $D_{\Theta}$  with weights  $\Theta_D$

Under certain conditions on  $f$ , finding  $\mathbb{W}$  is equivalent to solving

$$\max_{\Theta_D \in \mathcal{W}} \mathbb{E}_{x \sim p_{\text{data}}} D_{\Theta_D}(x) - \mathbb{E}_{x \sim p_{\text{model}}} D_{\Theta_D}(x)$$

where  $\mathcal{W}$  is a "compact" space of possible weights

Since the Discriminator no longer produces binary categorical values, it is more appropriate to call it a *Critic*.

That is: we solve for Critic weights such that the scores it produces have a large spread between Real and Fake.

## But: what does this mean ?

For those (like me) struggling with the math, here are the implications from a practical perspective

- Scores for true Real is negative, for Fake is positive
- $\mathcal{L}_G$  will implement: minimize (make most negative) the score assigned to Fakes
- $\mathcal{L}_D$  will implement: "maximize the spread of scores between Real and Fake"
  - by minimizing the sum of
    - sum of scores for Real examples
    - minus sum of scores for Fake examples (i.e., Discriminator goal is for Fakes to have positive scores)
- "Compact"  $\Theta_D$  will be achieved by clipping
  - restricting elements of  $\Theta_D$  to a small numerical range
  - by clipping the weights after a gradient update step for  $D$

- $\mathcal{L}_D$  will dispense with the log since the Discriminator produces scores rather than probabilities
  - we see terms  $D(\mathbf{x}^{(i)})$  and  $1 - D(\mathbf{x}^{(i)})$
  - rather than  $\log D(\mathbf{x}^{(i)})$  and  $1 - \log D(\mathbf{x}^{(i)})$

When we visit the code, we will see these elements in practice.

## Did I really need to change to EMD ?

The Wasserstein GAN avoids many of the problems associated with the plain GAN.

To some extent, this is due to replacing the Discriminator with a Critic

- unbounded scores in the WGAN versus bounded probabilities in the plain GAN

- There are mathematical problems with Expectation Maximization (KL distance) and Jensen-Shannon (JS) distance
  - the terms  $\log(p_{\text{model}}(\mathbf{x}))$  and  $\log(p_{\text{data}}(\mathbf{x}))$  appear
  - if  $p_{\text{data}}$  and  $p_{\text{model}}$  don't completely overlap (a possibility especially early in training)
    - we take logs of 0, which is infinite (negative)
  - No such problem with EMD

- No vanishing gradients with EMD
  - with KL and JS distance the true derivative goes to 0
  - no such problem with EMD
    - The Critic's scores are not bounded, so can't saturate
  - Thus: we can train the Discriminator to convergence immediately
    - No danger of being too good too soon
    - When we see code we will observe
      - The number of steps of Discriminator update is a multiple of the number of steps of Generator update
- No mode collapse with EMD
  - with a fixed Discriminator (classifier), the Generator in a plain GAN will seek out examples with highest probability of being mis-classified as Real



**Code**

